

Same 24 h diurnal α through both deployments — DDA's gain tracks how badly a uniform mapping wastes the fabric, not model size
(shaded = DDA worse on TTFT; 13B wins throughout, 7B only below ~ 0.6 load, 70B's Pool L is too small to serve the latency class at all)

