


Static Hybrid TP-PP: 32 CXL Devices, TP=16, PP=2

Llama2-70B | 80 Decoder Layers | Embedding dim = 8192



Config Summary

TP = 16 devices per pipeline stage

PP = 2 pipeline stages total

32 devices 16 × 2 = total

FC layers → sharded across 16 devs

Norm / Attn → single master device

Tokens in-flight 2 (one per PP stage)

TTFT ≈ 2 × stage latency

CXL comm. All-Reduce per FC layer

Why PP=2?

While Stage 2 processes
Token A, Stage 1 starts
Token B in parallel

→ 2× throughput vs PP=1